

Identifying Fraud from Enron Emails and Financial Data

Introduction

In 2000, Enron was one of the largest companies in the United States. By 2002, it had collapsed into bankruptcy due to widespread corporate fraud. In the resulting Federal investigation, there was a significant amount of typically confidential information entered into the public record, including tens of thousands of emails and detailed financial data for to executives.

Utilizing the classifiers and techniques taught in Into to Machine Learning Class, I built a classifier to detect if a person is culpable or not.

Short Questions

Summarize for us the goal of this project and how machine learning is useful in trying to accomplish it. As part of your answer, give some background on the dataset and how it can be used to answer the project question. Were there any outliers in the data when you got it, and how did you handle those?

The features in the data fall into three major types, namely financial features, email features and POI labels.

- financial features: salary, deferral_payments, total_payments, loan_advances, bonus, restricted_stock_deferred, deferred_income, total_stock_value, expenses, exercised_stock_options, other, long_term_incentive, restricted_stock, director_fees
- email features: to_messages, email_address, from_poi_to_this_person, from_messages, from_this_person_to_poi, shared_receipt_with_poi
- poi label(A total of 18 entries was labelled as POI)

The goal of this project was simply to leverage the features above in order to mark an individual as a person of interest.

In order to detect outlier and explore the data, I first built a code to convert the dictionary to csv file and is available with the submission. Later, I explored all the values to find out which features had a lot of values and which did not. Immediately looking at the csv, I noticed that the number of data is only 146 and so individually scanning each name - I found the following data points redundant

- Total: Does not convey information pertaining to any individual, hence marked as outlier.
- THE TRAVEL AGENCY IN THE PARK: It does not represent any individual and hence was removed.
- LOCKHART EUGENE E: This record contained no useful data.

After cleaning the data only 143 records remained.

What features did you end up using in your POI identifier, and what selection process did you use to pick them? Did you have to do any scaling? Why or why not? As part of the assignment, you should attempt to engineer your own feature that doesn't come ready-made in the dataset—explain what feature you tried to make, and the rationale behind it. If you used an algorithm like a decision tree, please also give the feature importances of the features that you use.

To pick best features, I used the select k best features of scikit learn. After trying different K values, I decided to go with K value as 8 as it gave the best performance. Folling is the list of top 8 features I selected. The K-best approach is an automated univariate feature selection algorithm, and in using it.

```
'exercised_stock_options' : 24.81
'total_stock_value' : 24.18
'bonus' : 20.79
'salary' - 18.28
'deferred_income' : 11.45
'long_term_incentive' : 9.92
'restricted_stock' : 9.21
'total_payments' - 8.77
```

These features did not include any email features. Therefore I decided to go with `shared_receipt_with_poi`. In order to include two more features that would represent an aggregate value for both financial and email features, I decided to add two more features, they are as follows:

- `email_interaction_ratio`: which was a ratio of the total number of emails to and from a POI to the total emails sent or received.
- `financial_aggregate`: This was the combined sum of `exercised_stock_options`, `salary`, and `total_stock_value`. This captured both liquid and semi solid wealth an individual has.

I scaled all features using a min-max scaler. This ensures features are evenly balanced and overcomes the disparity due to the units of financial and email features.

Thus after this step we have a total of 11 features to go ahead with.

What algorithm did you end up using? What other one(s) did you try?

After having performed various ml related projects during my undergraduate studies, a two class problem usually is best for logistic regression and K-means clustering. K-means clustering with PCA and mahalanobis distance provides a very fortified technique for two class detection. However, I went with logistic regression as my final algorithm.

I tried several algorithms, with a K-means clustering algorithm performing reasonably sufficient. I also tested a support vector machine, a random forest classifier, and stochastic gradient descent. The best results I got were from logistic regressor.

following were the parameters I tuned: - Logistic regression: `C` (inverse regularization parameter), `tol` (tolerance), and `class_weight` (over/undersampling) - K-means clustering: `tol`

K-means clustering was initialized with `K` (`n_clusters`) of 2 to represent POI and non-POI clusters. It performed well with the 11 set of features.

Auto-weighting in the case of logistic regression caused a dip in precision and hence I decided to use evenly balanced features.

The other algorithms were tuned experimentally, with unremarkable improvement.

What is validation, and what's a classic mistake you can make if you do it wrong? How did you validate your analysis?

Validation is performed to ensure that a machine learning algorithm generalizes well. The classic problem that can occur is over-fitting. This happens when we overfit the training data and perform really well in it due to which there is a considerable dip in the performance in the other two datasets(cross validation and testing dataset).

I validated my result using two techniques - bootstrapping (`cleaner.py`) - k fold method

Give at least 2 evaluation metrics, and your average performance for each of them. Explain an interpretation of your metrics that says something human-understandable about your algorithm's performance.

I am using precision and recall as the evaluation metric. - Precision: The ratio of true positives to the records that are actually POIs. This describes the occurrence of false alarms - Recall: captures the ratio of true positives to the records flagged as POIs, which describes sensitivity.

I did not choose accuracy because with such a small dataset using it would mean that a simple heuristic with all data marked false would give an accuracy of more than 85%

Validation 1 (Stratified K-folds, $K=3$)

Classifier	Precision	Recall	Features
Logistic Regression	0.422	0.170	11
K-means Clustering, $K=2$	0.176	0.062	11

Validation 2 randomised sampling($n=1000$)

Classifier	Precision	Recall	Features
Logistic Regression	0.329	0.209	11
K-means Clustering, $K=2$	0.315	0.375	11

Both algorithms do well inspite of the dataset being noisy. Let us understand what the values are actually speaking to us. Now, in terms of precision and recall you would want a high recall in such a case. Simply, because you want to be suspecting people. Having a high precision would mean that we are looking for too strict of a conditions to flag an individual. Hence, recall plays a high role in such a case.

Conclusion

The dataset was sparse and most algorithms will perform well only if given a decent number of data to learn. Had there been a large dataset random forest would also work very well. However, I am more interested in what a classic anomaly detection system would do in such a case. A simple combination of - K-means - PCA - Mahalanobis distance measure would result in a very robust outlier detection algorithm and probably would have both a high recall and precision value.

Resources and References

- [Introduction to Machine Learning \(Udacity\)](#)
- [Machine Learning \(Stanford/Coursera\)](#)
- [scikit-learn Documentation](#)